

lang_identification

November 29, 2023

1 Scraping YouTube for Language Identification and Toxicity Detection Tasks

1.1 Assignment #2

1.1.1 Practical Data Science course, MSc in Data Science (2023/2024)

Dimitris Tsirmpas MSc in Data Science f3352315 Athens University of Economics and Business

In this project we attempt to achieve the following goals:

- Creating a language dataset including Greeklish
- Crawling YouTube videos which include both Greek and Greeklish comments
- Training a language identification classifier
- Training a LLM-based toxicity classifier
- Using the LLM classifier to produce data for, and train a traditional ML toxicity classifier
- Applying our language identification and toxicity classifiers on the crawled YouTube videos and identifying interesting facts and trends

1.2 Directory Structure

The project is structured as follows:

Main files: - lang_identification.ipynb: is the main Jupyter Notebook containing the project code
- prompts.pdf: Supplemental material containing the prompts used for the toxicity LLM classifier
- report.pdf: Supplemental material containing Figures, Tables and analysis on the results of the project

Directories: - src: a library of general functions for Data Science tasks - tasks: task-specific modules
- data: the input data - output: the output data (.csv) - results: Graphs, Tables and Figures produced in the project

1.3 Disclaimers

- Most documentation was generated by ChatGPT, and was manually corrected / augmented where necessary
- In each case where code has been obtained from outside sources, is clearly listed either in comments or in the markdown explaining the code block

- Most code implementation is “hidden” in the `src` and `tasks` modules, as stated in the section above, which contain all the documentation and implementation details

```
[1]: from time import time

start = time()
```

Before we begin, we need to have a baseline rule-based model which we will originally use to classify our input data. We thus create a model which predicts a text’s language only using regex rules. This classifier obviously does not need to be trained (fitted), and thus we will only use it for predictions. The implementation of our model was inspired by [this excellent LinkedIn post](#).

We will define our language identification task as identifying one of three languages: Greek, English and Greeklish. We will also include an “other” category for all other languages.

1.4 Regex classifier

```
[2]: from sklearn.base import BaseEstimator, RegressorMixin
import numpy as np
import re

class RegexClassifier(BaseEstimator, RegressorMixin):
    """
    Language Classifier using Regular Expressions
    """

    language_regex_dict = {
        'el': r'([\u0370-\u03ff\u1f00-\u1fff]+\s?', # Greek
        'en': r'([a-zA-Z]+\s?', # English
    }

    # STATIC INITIALIZATION BLOCK
    # pre-compile all expressions to save execution time
    for lang in language_regex_dict.keys():
        language_regex_dict[lang] = re.compile(language_regex_dict[lang])

    def __init__(self, include_greeklish=True):
        """
        Initialize a new :class:`RegexClassifier` instance.

        :param include_greeklish: Optional. If True, include Greeklish in the
        ↪ language identification process.
        """
        self.include_greeklish = include_greeklish

    def predict(self, x):
```

```

preds = []

for text in x:
    pred = "other"

    if self.include_greeklish and \
        re.search(RegexClassifier.language_regex_dict["el"], text) and \
        re.search(RegexClassifier.language_regex_dict["en"], text):
        pred = "greeklish"
    else:
        for lang_code, regex_pattern in RegexClassifier.
↪language_regex_dict.items():
            match = re.search(regex_pattern, text)
            if match:
                pred = lang_code
                break
        preds.append(pred)
return np.array(preds)

def fit(self, x, y):
    return self

```

```

[3]: # Example usage
user_input_text = ["", "this is an english sentence", ""]
detected_language = RegexClassifier().predict(user_input_text)
print(f"Detected Language Code: {detected_language}")

```

Detected Language Code: ['other' 'en' 'el']

1.5 Defining the Gold Standard

The first task we ought to complete is derive our “gold” dataset, aka the dataset on which our classifiers will be trained on.

1.5.1 Greek-English identification dataset

We will begin by creating a base dataset including three of the four categories of data we defined in our language identification task (those being “Greek”, “English” and “Other”).

For this we will be using a subset of the [papluca language identification dataset](#), which is available to us through [huggingface](#).

We download the dataset and turn it into a pandas Dataframe, so we can easily combine it with other sources later on.

```

[4]: import pandas as pd

def dataset_to_pd(dataset_dict: dict) -> pd.DataFrame:

```

```

"""
Convert a dictionary of datasets into a single pandas DataFrame.

The datasets are assumed to be stored as values in the input dictionary,
with corresponding labels as keys.

:param dataset_dict: A dictionary where keys are labels and values are
↳datasets.
:type dataset_dict: dict

:return: A concatenated pandas DataFrame with an additional 'set' column
        indicating the original label of each row.
:rtype: pd.DataFrame

:Example:

.. code-block:: python

>>> dataset_dict = {'A': dataset_A, 'B': dataset_B, 'C': dataset_C}
>>> result_df = dataset_to_pd(dataset_dict)
>>> print(result_df)
      set ... (columns of the datasets)
0    A   ...
1    A   ...
...   ...
n    C   ...
"""
df_ls = []
label_array = np.empty(shape=(sum([len(dataset) for dataset in dataset_dict.
↳values()]), dtype=object)
last_idx = -1

for label, dataset in dataset_dict.items():
    new_last_idx = len(dataset) + last_idx
    label_array.put(np.arange(last_idx+1, new_last_idx+1, 1), label)
    last_idx = new_last_idx

    df_ls.append(pd.DataFrame(dataset))

full_df = pd.concat(df_ls, ignore_index=True)
full_df["set"] = label_array
full_df.insert(0, "set", full_df.pop("set"))

return full_df

```

```
[5]: from datasets import load_dataset
```

```
dataset_dict = load_dataset("papluca/language-identification")
dataset_dict
```

```
[5]: DatasetDict({
      train: Dataset({
            features: ['labels', 'text'],
            num_rows: 70000
        })
      validation: Dataset({
            features: ['labels', 'text'],
            num_rows: 10000
        })
      test: Dataset({
            features: ['labels', 'text'],
            num_rows: 10000
        })
  })
```

```
[6]: lang_df = dataset_to_pd(dataset_dict)
      lang_df
```

```
[6]:      set labels      text
0      train      pt  os chefes de defesa da estónia, letónia, lituâ...
1      train      bg
2      train      zh
3      train      th      honeychurch ...
4      train      ru
...      ...      ...
89995     test      zh
89996     test      tr  Örneğin, teşhis Yunanca bir kelimedenden alındı (...
89997     test      vi  Nếu lite/light chỉ đơn giản là mô tả một đặc t...
89998     test      bg      ,      ,      ...
89999     test      pl      Mam dla ciebie kilka propozycji:
```

[90000 rows x 3 columns]

```
[7]: en_gr_cond = lang_df.labels.eq("el") | lang_df.labels.eq("en")
      en_gr_df = lang_df.loc[en_gr_cond, ["labels", "text"]]
      en_gr_df
```

```
[7]:      labels      text
18      el      Π      ,      ...
39      en      Didnt really seem to work much.
40      el      A      ...
49      en  Highly recommend for those who don't like bein...
75      el      E      .
```

```

...
89961    en  It's super cute, really soft. Print is fine bu...
89965    en  One of them worked, the other one didn't. Ther...
89978    en  I only received one out of the three strikers :(
89982    el  0
89986    el  T Abeam      Arabella,

```

[9000 rows x 2 columns]

We sample a set of 2000 data records belonging to other languages and tag them as “other”. The exact number is arbitrary, although in general we would like to have a size: - large enough to contain most common words from foreign languages - small enough to not unnecessarily burden training or bias our classifiers towards rare occurrences in our operational set (the dataset which we claim little to no prior knowledge)

```

[8]: others_df = lang_df.loc[~en_gr_cond, ["labels", "text"]]
      others_df = others_df.sample(2000)
      others_df.labels = "other"
      others_df

```

```

[8]:      labels      text
18882  other
54679  other      2016 6  18
2154   other  Il Presidente del Consiglio dei Ministri Mario...
32815  other  1946 ' de ingiliz görev personeli tarafından i...
24227  other  Ugoda zawiera 4,1 miliona dolarów na honoraria...
...
12579  other
38271  other      100
24788  other
37396  other      linguist      weinreich
65858  other  l'agenzia internazionale per l'energia atomica...

```

[2000 rows x 2 columns]

```

[9]: gold1_df = pd.concat([en_gr_df, others_df], axis=0, ignore_index=True,
      ↪copy=False)
      gold1_df

```

```

[9]:      labels      text
0       el  Π
1       en      Didnt really seem to work much.
2       el  A
3       en  Highly recommend for those who don't like bein...
4       el      E
...
10995  other
10996  other      100

```

```

10997 other          '
10998 other          linguist      weinreich      ...
10999 other  l'agenzia internazionale per l'energia atomica..

```

```
[11000 rows x 2 columns]
```

1.5.2 Greek-Greeklsh identification dataset

A much harder task will be to include Greeklsh in our dataset. There are only few papers dedicated to Greeklsh language identification or translation [1,2]. Out of these, only one provides a [comprehensive Greeklsh dataset](#) [2], which is not publically available.

[1] Aimilios Chalamandaris, Athanassios Protopapas, Pirros Tsiakoulis, and Spyros Raptis. 2006.

[2] A. Chalamandaris, P. Tsiakoulis, S. Raptis, G. Giannopoulos, and G. Carayannis. 2004. Bypa

We thus need to create our own dataset. We begin with crawling a Greek gaming forum, which we have verified contains many posts in Greeklsh and almost none in English.

```
[10]: head_url = "https://forum.warmane.com"
      warmane_url = "https://forum.warmane.com/forumdisplay.php?f=20"
```

```
[11]: from src.crawling import fetch_soup
      from tasks.warmane import parse_warmane_thread
      from tqdm import tqdm

      threads = []

      for page in range(1, 9):
          url = warmane_url + f"&page={page}"
          soup = fetch_soup(url)

          print(f"Processing page {page} of 8...")
          thread_tags = soup.find_all("li", {"class": "threadbit"})
          for thread_tag in tqdm(thread_tags):
              thread = parse_warmane_thread(head_url, thread_tag)
              threads.append(thread)
```

Processing page 1 of 8...

```

100%|
  | 20/20 [00:04<00:00, 4.15it/s]

```

Processing page 2 of 8...

```

100%|
  | 20/20 [00:03<00:00, 5.18it/s]

```

Processing page 3 of 8...

```

100%|
  | 20/20 [00:04<00:00, 4.97it/s]

```

Processing page 4 of 8...

100%|
| 20/20 [00:03<00:00, 5.39it/s]

Processing page 5 of 8...

100%|
| 20/20 [00:03<00:00, 5.05it/s]

Processing page 6 of 8...

100%|
| 20/20 [00:03<00:00, 5.71it/s]

Processing page 7 of 8...

70%|
| 14/20 [00:02<00:01, 4.97it/s]

ERROR: Failed to get information on post
<https://forum.warmane.com/showthread.php?t=272585>

100%|
| 20/20 [00:03<00:00, 5.35it/s]

Processing page 8 of 8...

78%|
| 7/9 [00:01<00:00, 5.45it/s]

ERROR: Failed to get information on post
<https://forum.warmane.com/showthread.php?t=278731>

100%|
| 9/9 [00:01<00:00, 5.66it/s]

```
[12]: import itertools

# flatten nested lists
posts = set(itertools.chain.from_iterable([thread.posts for thread in threads]))
len(posts)
```

[12]: 415

```
[13]: import pandas as pd

warmane_df = pd.DataFrame.from_records([post.__dict__ for post in posts],
    ↪ index="id")
warmane_df.reply_to = warmane_df.reply_to.fillna(-1).astype(int)
warmane_df
```

```
[13]:      thread_id      author \
id
```


2926596	384475	Ripsin
2473988	300013	v4gflo
2420747	290921	AlexPan
2981903	399822	xAchillesGate4x
2879517	371804	Csdas
...
2877428	353812	Shiverbro
3069941	423611	crystallenia898
2801654	350071	Draculation
2873339	370241	Ripsin
2410495	289030	boolouk

	contents	date \
id		
2926596	Kalhspera paides,\n\r\nEimai arketo kairo ston...	2018-05-22
2473988	geia sas.psaxnw ellhniko guild ston Deathwing ...	2015-06-17
2420747	K , . \...	2015-03-24
2981903	K . Ψ E active raidin...	2019-03-03
2879517	Opoios gnwrizei kati as mou kanei /w Dremoria ...	2017-11-29
...
2877428	kalos private aksizei na ksekiniseis paidia?	2017-11-21
3069941	E ...	2020-07-26
2801654	Bump! ICC25 6/12	2017-05-07
2873339	Kalhspera tha ithela na rwthsw an kapoios gnwr...	2017-11-07
2410495	E , ...	2015-03-13

	reply_to
id	
2926596	-1
2473988	-1
2420747	-1
2981903	-1
2879517	-1
...	...
2877428	2875915
3069941	3068345
2801654	2795443
2873339	-1
2410495	2409274

[415 rows x 5 columns]

We will also clear out empty posts.

```
[14]: empty_contents = warmane_df.contents.apply(lambda x: x.isspace() | len(x)==0)
      warmane_df[empty_contents]
```

```
[14]:
```

	thread_id	author	contents	date	reply_to
id					
3082464	427259	malakas17		2020-10-20	3081822
3113236	427259	malakas17		2021-05-12	3113009
3099161	431660	malakas17		2021-02-10	3096432
3113819	427259	malakas17		2021-05-16	3113236
3099593	427259	boonick		2021-02-14	3093400
3081820	427259	malakas17		2020-10-16	3080427
3081822	427259	malakas17		2020-10-16	3081820

```
[15]: warmane_df = warmane_df[~empty_contents]
```

While this dataset fits our needs, it is by no means large enough to accurately model Greeklish on traditional NLP models.

Thus, we turn our attention early to YouTube scraping. YouTube is one of the few sites featuring lively comment sections in Greek informal enough for Greeklish to be present. Knowing that Greeklish are generally more prevalent towards younger generations, we can select certain videos where this demographic is present to scrape for comments. In our case, we select 5 Greek gaming videos.

```
[16]: from src.crawling import ChromeDriverManager, jupyter_options
```

```
ChromeDriverManager.set_options(jupyter_options())
```

```
[17]: from tasks.youtube import extract_search_results, extract_comments, \
      ↪scrape_youtube

greek_yt_urls = ["https://www.youtube.com/watch?v=4Y2gxkqbsbA",
                  "https://www.youtube.com/watch?v=31LcJ9gqQvA",
                  "https://www.youtube.com/watch?v=1cZXAQ1JEJo",
                  "https://www.youtube.com/watch?v=x7lnS6jMS64",
                  "https://www.youtube.com/watch?v=ImilczGN-00"]
scrape_results = []

for url in tqdm(greek_yt_urls):
    scrape_results.append(scrape_youtube(ChromeDriverManager.get(), url, \
    ↪max_scrolls=10, verbose=False))
```

```
0%|
| 0/5 [00:00<?, ?it/s]

Creating new driver...
New driver online.

100%|
| 5/5 [01:35<00:00, 19.00s/it]
```

```
[18]: from tasks.youtube import extract_comments
```

```
all_comments = []
for result in scrape_results:
    comments, _ = extract_comments(result)
    all_comments += comments

all_comments = pd.Series(all_comments)
```

We now combine our two Greeklish datasets:

```
[19]: greeklsh_series = pd.concat([warmane_df.contents, all_comments])
greeklsh_series
```

```
[19]: 2926596    Kalhspera paides,\n\r\nEimai arketo kairo ston...
      2473988    geia sas.psaxnw ellhniko guild ston Deathwing ...
      2420747    K          ,          . \...
      2981903    K          . Ψ    E    active raidin...
      2879517    Opoios gnwrizei kati as mou kanei /w Dremoria ...

      ...

      744                                     Π
      745                                First of all
      746                                First
      747                                     .
      748                                     Π
      Length: 1157, dtype: object
```

We filter out empty and “junk” comments.

```
[20]: conditions = (greeklsh_series.apply(lambda x: len(x) != 0)) & \
              (greeklsh_series.apply(lambda x: "RRR" not in x)) & \
              (greeklsh_series.apply(lambda x: "PPP" not in x)) & \
              (greeklsh_series.apply(lambda x: "First" not in x))
cleared_greeklsh_series = greeklsh_series[conditions]
cleared_greeklsh_series
```

```
[20]: 2926596    Kalhspera paides,\n\r\nEimai arketo kairo ston...
      2473988    geia sas.psaxnw ellhniko guild ston Deathwing ...
      2420747    K          ,          . \...
      2981903    K          . Ψ    E    active raidin...
      2879517    Opoios gnwrizei kati as mou kanei /w Dremoria ...

      ...

      742                                     Π
      743                                Gianni    Pubg
      744                                     Π
      747                                     .
      748                                     Π
```

Length: 1139, dtype: object

And annotate the entire Greeklish dataset using our rules-based (Regex) classifier.

Since we selected videos exclusively in Greek, we can safely assume that the vast majority of comments not featuring Greek characters are in Greeklish. We can also safely assume from prior knowledge that Greeklish comments will conversly not feature any Greek characters.

We thus classify all comments with English charactes are Greeklish.

```
[21]: regex_model = RegexClassifier(include_greeklish=False)
      preds = regex_model.predict(cleared_greeklish_series)
```

We can briefly verify that our assumptions are correct:

```
[22]: cleared_greeklish_series[preds=="en"]
```

```
[22]: 2926596    Kalhspera paides,\n\r\nEimai arketo kairo ston...
      2473988    geia sas.psaxnw ellhniko guild ston Deathwing ...
      2879517    Opoios gnwrizei kati as mou kanei /w Dremoria ...
      2959390           Bubblethesap Icecrown wotlk horde belf
      2947119    den se vrisko kane add evvi  .\nmou leei den u...

      ...

      730                                           Geia
      731                                           Lol
      732    Ante Pali me ta atoma pou einai first... Mhn ...
      733           Protos Protos molis vgike
      740                                           Hafa
      Length: 430, dtype: object
```

```
[23]: cleared_greeklish_series[preds=="el"]
```

```
[23]: 2420747    K          ,          . \...
      2981903    K          . Ψ    E    active raidin...
      2959391    K          (properties) ...
      2719776    Originally Posted by celphecil\n\nK    Σ ...
      2971700    E    guild ,          runs ICC10...

      ...

      741                                           Π    like
      742                                           Π
      743           Gianni    Pubg
      744                                           Π
      748                                           Π
      Length: 703, dtype: object
```

```
[24]: labels = np.where(preds=="en", "greeklish", "el")
      gold2_df = pd.DataFrame({"labels": labels, "text": cleared_greeklish_series})
      gold2_df
```

```
[24]:
```

	labels	text
2926596	greeklish	Kalhspēra paidēs,\n\r\nEimai arketo kairo ston...
2473988	greeklish	geia sas.psaxnw ellhniko guild ston Deathwing ...
2420747	el	K , . \...
2981903	el	K . Ψ E active raidin...
2879517	greeklish	Opoios gnwrizei kati as mou kanei /w Dremoria ...
...
742	el	Π
743	el	Gianni Pubg
744	el	Π
747	el	.
748	el	Π

[1139 rows x 2 columns]

Having our Greek-English-Other and our Greek-Greeklish datasets we can now combine them to form our gold dataset.

```
[25]: gold_df = pd.concat([gold1_df, gold2_df])
gold_df
```

```
[25]:
```

	labels	text
0	el Π	,
1	en	Didnt really seem to work much.
2	el A	...
3	en	Highly recommend for those who don't like bein...
4	el	E .
...
742	el	Π
743	el	Gianni Pubg
744	el	Π
747	el	.
748	el	Π

[12139 rows x 2 columns]

```
[26]: import os

OUTPUT_DIR = "output"

def csv_output(df: pd.DataFrame, filename: str) -> None:
    """
    Save a pandas DataFrame to a CSV file.

    :param df: The DataFrame to be saved.
    :type df: pd.DataFrame
```

```

:param filename: The name of the CSV file.
:type filename: str

:return: This function does not return anything.
:rtype: None
"""

file = os.path.join(OUTPUT_DIR, filename)
df.to_csv(file, encoding = 'utf8')
print(f"File saved successfully as {file}")

```

```
[27]: csv_output(gold_df, "gold.csv")
```

File saved successfully as output\gold.csv

1.6 Youtube Crawling

Our search strategy consists of:

- Searching YouTube for a specific Greek topic
- Getting the links and video names from the search
- For each link, crawl the comments for a set number of scrolling actions

By using the YouTube search function we can guarantee that the crawled videos will be relevant (many comments in Greek-Greeklish) and popular (large number of comments).

We will repeat this procedure twice for two distinct groups of videos: greek songs, since their comments are usually in more formal Greek, and Greek gaming videos where, as discussed above, because of the demographic Greeklish are more prevalent.

```
[28]: from tasks.youtube import extract_search_results, extract_comments
```

```

# "greek songs" search in Greek
song_search_url = "https://www.youtube.com/results?
↳search_query=%CE%B5%CE%BB%CE%BB%CE%B7%CE%BD%CE%B9%CE%BA%CE%B1+%CF%84%CF%81%CE%B1%CE%B3%CE%B
search_soup = scrape_youtube(ChromeDriverManager.get(), song_search_url,
↳max_scrolls=5, verbose=True)
results_search_song = extract_search_results(search_soup)

```

```

Scrolling (0 out of max 5)...
Scrolling (1 out of max 5)...
Scrolling (2 out of max 5)...
Scrolling (3 out of max 5)...
Scrolling (4 out of max 5)...
Scrolling (5 out of max 5)...

```

```
[29]: gaming_search_url = "https://www.youtube.com/results?
↳search_query=greek+fortnite"
gaming_soup = scrape_youtube(ChromeDriverManager.get(), gaming_search_url,
↳max_scrolls=5, verbose=True)
```

```
results_search_gaming = extract_search_results(gaming_soup)
```

```
Scrolling (0 out of max 5)...
Scrolling (1 out of max 5)...
Scrolling (2 out of max 5)...
Scrolling (3 out of max 5)...
Scrolling (4 out of max 5)...
Scrolling (5 out of max 5)...
```

```
[30]: results_df = pd.DataFrame({"title": results_search_song[0] +
    ↪results_search_gaming[0],
    "link": results_search_song[1] +
    ↪results_search_gaming[1],
    "source": np.
    ↪array(len(results_search_song[0])*["song"] +
    ↪len(results_search_gaming[0])*["gaming"]) })
results_df
```

```
[30]:
```

	title \	link	source
0	\n\nGreek Hits 2023 Non-Stop Mix by Elegant ...	/watch?v=RcSAggke-_U&pp=ygUjzrX0u867zrf0vc65zr...	song
1	\n\n00's GREEK MIX KAPSOURA EDITION\n	/watch?v=isCeE38TrXA&pp=ygUjzrX0u867zrf0vc65zr...	song
2	\n\nM N.1 () - 100 ...	/watch?v=p5g82ta4sTk&pp=ygUjzrX0u867zrf0vc65zr...	song
3	\n\nThe Greek '90s Dance NonStopMix OFFICIAL...	/watch?v=GEPvs6JA_c&pp=ygUjzrX0u867zrf0vc65zr...	song
4	\n\nTA ΛΑΙΚΑ THE TABEPNAΣ NON STOP MIX - Π ...	/watch?v=C4f3xcZzr3s&pp=ygUjzrX0u867zrf0vc65zr...	song
..
262	\n\nH TEAETTAIA ΦΟΡΑ ΠΟΤ ΠΑΙΖΩ FORTNITE ONLY U...	/watch?v=KLBYBvxiTvQ&pp=ygU0Z3JlZWsgZm9ydG5pdG...	gaming
263	\n\n UNREAL STREAM Fortnite Greek Live Stream...	/watch?v=4lyI1hJlnKE&pp=ygU0Z3JlZWsgZm9ydG5pdG...	gaming
264	\n\nΟ ΧΑΜΕΝΟΣ ΤΡΟΕΙ ΜΙΑ ΑΠΑΙΣΙΑ PIZZA *ΜΗΑΙΑΧ*...	/watch?v=qFgEIRgj9pE&pp=ygU0Z3JlZWsgZm9ydG5pdG...	gaming
265	\n\nΔΟΚΙΜΑΖΩ ΤΟ CHAPTER 2 ΣΤΟ FORTNITE ONLY UP...	/watch?v=e7sawZ1bL6c&pp=ygU0Z3JlZWsgZm9ydG5pdG...	gaming
266	\n\nΑΤΤΟΕ ΕΧΕΙ '0Views'(React Σ Montages M ...	/watch?v=6ay2HZwz2sA&pp=ygU0Z3JlZWsgZm9ydG5pdG...	gaming

```
[267 rows x 3 columns]
```

```
[31]: results_df.title = results_df.title.apply(lambda x: x.strip())
results_df.link = results_df.link.apply(lambda x: "https://www.youtube.com" + x)
results_df
```

```
[31]:
```

	title \	link	source
0	Greek Hits 2023 Non-Stop Mix by Elegant Gree...	https://www.youtube.com/watch?v=RcSAggke-_U&pp...	song
1	OO's GREEK MIX KAPSOURA EDITION	https://www.youtube.com/watch?v=isCeE38TrXA&pp...	song
2	M N.1 () - 100 ...	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
3	The Greek '90s Dance NonStopMix OFFICIAL Part 1	https://www.youtube.com/watch?v=GEPvsn6JA_c&pp...	song
4	TA ΛΑΙΚΑ THE TABEPNAΣ NON STOP MIX - Π ...	https://www.youtube.com/watch?v=C4f3xcZzr3s&pp...	song
...
262	H TEΛΕΤΤΑΙΑ ΦΟΡΑ ΠΟΤ ΠΑΙΖΩ FORTNITE ONLY UP CH...	https://www.youtube.com/watch?v=KLBYBvxiTvQ&pp...	gaming
263	UNREAL STREAM Fortnite Greek Live Stream Now...	https://www.youtube.com/watch?v=4lyI1hJlnKE&pp...	gaming
264	O XAMENOS TPΩEI MIA APAIΣIA PIZZA *MHAIAX* (F...	https://www.youtube.com/watch?v=qFgEIRgj9pE&pp...	gaming
265	ΔΟΚΙΜΑΖΩ ΤΟ CHAPTER 2 ΣΤΟ FORTNITE ONLY UP * R...	https://www.youtube.com/watch?v=e7sawZ1bL6c&pp...	gaming
266	ΑΤΤΟΣ ΕΧΕΙ '0Views'(React Σ Montages M '0'V...	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming

[267 rows x 3 columns]

We will use our rules-based (Regex) classifier to once again filter the videos, so they only include titles in Greek.

```
[32]: greeklsh_model = RegexClassifier(include_greeklsh=True)
preds = greeklsh_model.predict(results_df.title)
gr_res_df = results_df[(preds != "en") & (preds != None)]
gr_res_df
```

```
[32]:
```

	title \
2	M N.1 () - 100 ...
4	TA ΛΑΙΚΑ THE TABEPNAΣ NON STOP MIX - Π ...
6	N Θ - A (AI Cover)
7	E disco, 80s & 90s (Non-stop Party Mix)...
8	B K - T E Vasilis...
...	...
261	H ΣΕΖΟΝ 4 ΕΙΝΑΙ ΕΔΩ!!! NEO MAP, BATTLE PAS...


```

262 H ΤΕΛΕΤΤΑΙΑ ΦΟΡΑ ΠΟΤ ΠΑΙΖΩ FORTNITE ONLY UP CH...
264 Ο ΧΑΜΕΝΟΣ ΤΡΩΕΙ ΜΙΑ ΑΠΑΙΣΙΑ PIZZA *ΜΠΛΙΑΧ* (F...
265 ΔΟΚΙΜΑΖΩ ΤΟ CHAPTER 2 ΣΤΟ FORTNITE ONLY UP * R...
266 ΑΤΤΟΣ ΕΧΕΙ '0Views'(React Σ Montages Μ '0'V...

```

	link	source
2	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
4	https://www.youtube.com/watch?v=C4f3xcZzr3s&pp...	song
6	https://www.youtube.com/watch?v=HILSv0QV_bc&pp...	song
7	https://www.youtube.com/watch?v=Y9rbyT0ZNq4&pp...	song
8	https://www.youtube.com/watch?v=mQkIg8Rg3m4&pp...	song
..
261	https://www.youtube.com/watch?v=SMVnHVPRm_Q&pp...	gaming
262	https://www.youtube.com/watch?v=KLBYBvxiTvQ&pp...	gaming
264	https://www.youtube.com/watch?v=qFgEIRgj9pE&pp...	gaming
265	https://www.youtube.com/watch?v=e7sawZ1bL6c&pp...	gaming
266	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming

[238 rows x 3 columns]

Below we begin the process of crawling, processing and packaging the crawled data into a unified dataframe:

```

[33]: from selenium.common.exceptions import JavascriptException
import bs4

def scrape(urls: list[str]) -> list[tuple[str, bs4.BeautifulSoup]]:
    """
    Scrape YouTube videos using the provided URLs.

    :param urls: A list of YouTube video URLs to scrape.
    :type urls: list[str]

    :return: A list of tuples where each tuple contains the original URL and
             the corresponding BeautifulSoup object containing the scraped data.
    :rtype: list[tuple[str, bs4.BeautifulSoup]]

    :Example:

    .. code-block:: python

        >>> scrape_results = scrape(['https://www.youtube.com/watch?
↪v=example1', 'https://www.youtube.com/watch?v=example2'])
        >>> print(scrape_results)
        [('https://www.youtube.com/watch?v=example1', <BeautifulSoup object>),
↪ ('https://www.youtube.com/watch?v=example2', <BeautifulSoup object>)]

```

```

"""
scrape_results = []

print("Scraping videos...")
for url in tqdm(urls):
    try:
        scrape_results.append((url, scrape_youtube(ChromeDriverManager.
↳get(), url, max_scrolls=10, scroll_wait_secs=1.3, verbose=False)))
    except JavaScriptException:
        continue
    except Exception as e:
        print(e)
        continue
return scrape_results

def process_scraped(scrape_results: list[tuple[str, bs4.BeautifulSoup]]) -> pd.
↳DataFrame:
    """
    Process the scraped YouTube video data to extract comments and dates.

    :param scrape_results: A list of tuples where each tuple contains the
↳original URL
                                and the corresponding BeautifulSoup object containing
↳the scraped data.
    :type scrape_results: list[tuple[str, bs4.BeautifulSoup]]

    :return: A pandas DataFrame containing the processed data with columns
↳'link', 'text', and 'date'.
    :rtype: pd.DataFrame

    :Example:

    .. code-block:: python

        >>> processed_data = process_scraped([('https://www.youtube.com/watch?
↳v=example1', <BeautifulSoup object>), ('https://www.youtube.com/watch?
↳v=example2', <BeautifulSoup object>)])
        >>> print(processed_data)

           link                                text    date
0  https://www.youtube.com/watch?v=example1  Comment 1  2023-01-01
1  https://www.youtube.com/watch?v=example1  Comment 2  2023-01-02
2  https://www.youtube.com/watch?v=example2  Comment 3  2023-01-03
    """
scraped_urls = []
comments = []

```

```

dates = []

print("Processing comments...")
print(type(scrape_results[0]))
for url, result in tqdm(scrape_results):
    if result is not None:
        new_comments, new_dates = extract_comments(result)
        comments += new_comments
        dates += new_dates
        scraped_urls += ([url] * len(new_comments))
return pd.DataFrame({"link": scraped_urls, "text": comments, "date": dates})

def filter_comments(df: pd.DataFrame) -> pd.DataFrame:
    preds = greeklsh_model.predict(df.text)
    mask = ((preds != "el") & (preds != "greeklsh"))
    return comments_df[mask]

```

We will crawl a maximum of 150 videos, using a random uniform mix of both video groups.

Initially multi-threading was considered but the processes proved too demanding on main memory and CPU resources. Thus, we run the crawl in a single thread.

```

[34]: max_videos = 150
urls = gr_res_df.link.sample(max_videos) if len(gr_res_df.link) > max_videos
else gr_res_df.link

scraped = scrape(urls)
comments_df = process_scraped(scraped)

comments_df.date = comments_df.date.apply(lambda x: x.date() if x is not None
else None)
comments_df = filter_comments(comments_df)
crawl_df = pd.merge(gr_res_df, comments_df, how="inner", on="link")

```

Scraping videos...

```

100%|
  | 150/150 [34:02<00:00, 13.62s/it]

```

Processing comments...

```

<class 'tuple'>

```

```

100%|
  | 135/135 [01:06<00:00, 2.02it/s]

```

We again clear the dataset of empty comments and other anomalies:

```

[35]: crawl_df = crawl_df.dropna()
crawl_df = crawl_df[~crawl_df.text.apply(lambda x: len(x.strip())==0)]

```

```
crawl_df
```

[35]:

```
                                title \
0      M      N .1 (      ) - 100  ...
1      M      N .1 (      ) - 100  ...
2      M      N .1 (      ) - 100  ...
3      M      N .1 (      ) - 100  ...
4      M      N .1 (      ) - 100  ...
...
2692  ATTOΣ EXEI '0Views'(React Σ Montages M '0'V...
2693  ATTOΣ EXEI '0Views'(React Σ Montages M '0'V...
2694  ATTOΣ EXEI '0Views'(React Σ Montages M '0'V...
2695  ATTOΣ EXEI '0Views'(React Σ Montages M '0'V...
2696  ATTOΣ EXEI '0Views'(React Σ Montages M '0'V...

                                link  source \
0      https://www.youtube.com/watch?v=p5g82ta4sTk&pp...  song
1      https://www.youtube.com/watch?v=p5g82ta4sTk&pp...  song
2      https://www.youtube.com/watch?v=p5g82ta4sTk&pp...  song
3      https://www.youtube.com/watch?v=p5g82ta4sTk&pp...  song
4      https://www.youtube.com/watch?v=p5g82ta4sTk&pp...  song
...
2692  https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...  gaming
2693  https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...  gaming
2694  https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...  gaming
2695  https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...  gaming
2696  https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...  gaming

                                text      date
0      ANTIGUAS CANCIONES DE GRECIA PAIS NATAL DE MIS...  2022-11-29
1      Încă o zi petrecută cu muzica voastră fantas...  2022-11-29
2      FelicităriSuperb \nSă fiți mereu bine \nMomen...  2022-11-29
3      Lovely collection! Thank you <3  2022-11-29
4      Beautiful thank you so many memories awesome ...  2022-11-29
...
2692      Copy by fantaros  2021-11-29
2693      Fist  2021-11-29
2694      Yui  2022-11-29
2695      Uihhii  2022-11-29
2696      Zzfitdt  2021-11-29
```

[2359 rows x 5 columns]

And close the driver we used for our crawling to relinquish the resources we obtained from the OS.

[36]: `ChromeDriverManager.quit()`

1.7 Language Identification

In this part, we create our optimal ML classifier for the language identification task we outlined in Part 1.

We will use the `gold.csv` dataset we built in Part 1 to create a train, test and validation split for our models.

We use stratified sampling, since our datasets are wildly unbalanced, especially in respect to Greeklish. This will ensure an adequate number of Greeklish posts will be included in all splits.

Given this imbalance in fact, we are tempted to use undersampling or oversampling in order to prevent our classifiers from simply ignoring the class. This however will likely not aid us significantly, given that most models (such as Logistic Regression [3]) are robust to class imbalance [4].

[3] Gary King and Langche Zeng. 2001. "Logistic Regression in Rare Events Data." *Political Analysis* 163

[4] Stephan Kolassa (<https://stats.stackexchange.com/users/1352/stephan-kolassa>), Are unbalanced

We will use a train-test-validation split of 70%-10%-20%. We don't expect a great need for validation since we will be using relatively simple ML models. Additionally, the difficulty presented by the small sample size of Greeklish comments means we must focus as much data as possible in making sure our classifiers can understand Greeklish in the first place.

```
[37]: from src.ml import train_test_val_split
import matplotlib.pyplot as plt


data_train, data_val, data_test = train_test_val_split(gold_df, train_ratio=0.
↪7, val_ratio=0.1, test_ratio=0.2,
                                                    random_state=42,
↪stratify_col="labels")
# code block from Ioannis Pavlopoulos
axes = pd.DataFrame({"train": data_train.labels.value_counts(),
                    "validation": data_val.labels.value_counts(),
                    "test": data_test.labels.value_counts()})
    .plot.pie(subplots=True,
              textprops={'fontsize': 8},
              autopct=f'%1.1f%%', # print percent% results
              legend=False)

axes[0].set_xlabel(data_train.shape[0])
axes[1].set_xlabel(data_val.shape[0])
axes[2].set_xlabel(data_test.shape[0])

axes[1].text(0, -2, 'Number of records', ha='center')

plt.title("Train-Test-Val split for the Language Identification Dataset")
plt.tight_layout(pad=2.0)
```

```
plt.show()
```



We will encode our data as TF-IDF vectors. Count vectors could also work for this specific problem, but the computational cost of TF-IDF is minimal compared to acquiring the data fitting our classifiers.

```
[38]: from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer().fit(data_train.text)
x_train = vectorizer.transform(data_train.text)
y_train = data_train.labels
x_val = vectorizer.transform(data_val.text)
y_val = data_val.labels
x_test = vectorizer.transform(data_test.text)
y_test = data_test.labels
```

The metric we will be using is Macro-F1 average.

- **F1** is a metric used to balance the need for making sure our classifications for a category are both correct (precision) and represent as many of the actual cases of the category as possible (recall).
- **Macro-F1** is the unweighted average of all F1 metrics for each class. We choose Macro F1 instead of a weighted average because
 - We have an unbalanced dataset (Greeklish data are a small fraction of overall data)
 - We are much more interested in the small classes (here Greeklish)

Thus, we want to use a metric which favors both thorough and precise classifiers, and which also assigns equal importance to our smaller classes.

```
[39]: from sklearn.model_selection import cross_val_score

def cross_val_res(model, x, y, scoring=None, cv=10):
    """
    Minor utility method, wrapping cross_val_score.
    """
    if scoring is None:
```

```

        scoring = "f1_macro"
    res = cross_val_score(model, x, y, cv=cv, scoring=scoring)
    return res

```

```

[40]: from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import warnings

def get_statistics(y_test, y_pred):
    """
    Minor utility method printing average Macro F1 score and classification_
    ↪report
    as well as displaying the classifier's Confusion Matrix.
    """
    with warnings.catch_warnings():
        warnings.simplefilter("ignore")

        print(f"Macro F1: {f1_score(y_test, y_pred, average='macro',
    ↪zero_division=0)}")
        print(classification_report(y_test, y_pred, zero_division=0))
        ConfusionMatrixDisplay.from_predictions(y_test,
                                                y_pred,
                                                colorbar=True)

    plt.show()

```

For the Language Identification task the following models are considered:

- The previously implemented rules-based (Regex) model
- Naive Bayes
- Logistic Regression
- Random Forest
- Adaboost Model

1.7.1 Dummy Classifier

We will first run a “fake” classifier which only guesses the majority category.

This dummy model thus completely disregards the input features and serves as a useful baseline with which to compare the subsequent classifiers.

```

[41]: from sklearn.dummy import DummyClassifier


majority = DummyClassifier(strategy="most_frequent")
majority.fit(x_train, y_train)
majority_res = majority.predict(x_test)

```

```
get_statistics(y_test, majority_res)
```

Macro F1: 0.14709302325581394

	precision	recall	f1-score	support
el	0.42	1.00	0.59	1012
en	0.00	0.00	0.00	942
greeklish	0.00	0.00	0.00	79
other	0.00	0.00	0.00	395
accuracy			0.42	2428
macro avg	0.10	0.25	0.15	2428
weighted avg	0.17	0.42	0.25	2428



1.7.2 Regex Classification

Let's now compare our rules-based classifier with the baseline. This classifier does use the input features to make decisions, but in a very simple and naive way. It also does not benefit from any information that could be gained through training.


```
[42]: with warnings.catch_warnings():
      warnings.simplefilter("ignore")


      regex_model = RegexClassifier()
      res = cross_val_res(regex_model, data_train.text, y_train,
      ↪scoring="f1_macro")
      print(f"Regex Classifier mean macro F1: {res[0]:.4f}, std: {res[1]:.4f}")
```

Regex Classifier mean macro F1: 0.5690, std: 0.5614

```
[43]: regex_model = RegexClassifier(include_greeklish=True).fit(data_train.text,
      ↪y_train)
      regex_res = regex_model.predict(data_test.text)
      get_statistics(y_test, regex_res)
```

Macro F1: 0.5427840052990762

	precision	recall	f1-score	support
el	1.00	0.69	0.82	1012
en	0.74	1.00	0.85	942
greeklish	0.00	0.00	0.00	79
other	0.98	0.34	0.51	395
accuracy			0.73	2428
macro avg	0.68	0.51	0.54	2428
weighted avg	0.86	0.73	0.75	2428



We notice that this classifiers performs relatively well at classifying Greek and English, but is generally easily confused and can not catch Greeklish at all.

1.7.3 Naive Bayes

Naive Bayes is a very cheap and easy-to-interpret classifier, which checks for the probability that each individual word in the text will belong in any language. We generally want to use the simplest model for the job, and so we start with this reliable model which has proven itself in many fields in the past.

The `sklearn` library gives us access to many variations of Naive Bayes, each specialized in its own field. For this NLP task, we will be using `MultinomialNB`, which was suggested by [this blogpost](#).

```
[44]: from sklearn.naive_bayes import MultinomialNB

# naive bayes needs dense arrays to work
naive_x_train = x_train.toarray()
naive_x_test = x_test.toarray()

naive_model = MultinomialNB()
res = cross_val_res(naive_model, naive_x_train, y_train, cv=5)
print(f"Naive Bayes mean macro F1-score {res[0]:.4f}, std: {res[1]:.4f}")
```


Naive Bayes mean macro F1-score 0.7848, std: 0.8102

```
[45]: naive_model = MultinomialNB().fit(naive_x_train, y_train)
naive_res = naive_model.predict(naive_x_test)

get_statistics(y_test, naive_res)
```

Macro F1: 0.8651470835475744

	precision	recall	f1-score	support
el	0.92	1.00	0.96	1012
en	0.96	1.00	0.98	942
greeklish	1.00	0.47	0.64	79
other	0.99	0.80	0.88	395
accuracy			0.95	2428
macro avg	0.97	0.82	0.87	2428
weighted avg	0.95	0.95	0.94	2428



Compared to our rules-based classifier, this is a great step-up. Classification of Greek and English is very good, and the “other” languages are reliably identified.

Greeklish however are not reliably caught. The classification report states that when the classifier guesses Greeklish, it is always correct, but most Greeklish comments are confused with either English or Greek. Thus, the problem of distinguishing these three categories probably requires a more complex model.

1.7.4 Logistic Regression

LogisticRegression despite its name is a linear classifier, meaning that it attempts to linearly separate the data into distinct categories. This interpretation does not apply well to a NLP task, but means that the classifier retains some very useful properties:

- The solution we get is a global optimum, meaning that it's the best we can get with the provided data. This means no hyper-parameter tuning is necessary and we can use the classifier as-is.
- It's a simple and very easy to compute classifier, since it solves a (mathematically simple) linear problem, albeit with some restrictions (technically those restrictions force it to use gradient descent, but the calculations are much easier than say, a neural network)

```
[46]: from sklearn.linear_model import LogisticRegression

with warnings.catch_warnings():
    # ignore warnings about deprecated methods in libraries
    warnings.simplefilter("ignore")

    lr = LogisticRegression(max_iter=1000)
    res = cross_val_res(lr, x_train, y_train)
    print(f"Logistic Regression mean macro F1-score {res[0]:.4f}, std: {res[1]:.4f}")
```

Logistic Regression mean macro F1-score 0.9042, std: 0.8825


```
[47]: with warnings.catch_warnings():
    warnings.simplefilter("ignore")
    lr = LogisticRegression(max_iter=1000).fit(x_train, y_train)
    lr_res = lr.predict(x_test)

get_statistics(y_test, lr_res)
```

Macro F1: 0.9243052241829711

	precision	recall	f1-score	support
el	0.99	0.98	0.99	1012
en	0.99	0.99	0.99	942
greeklish	0.98	0.63	0.77	79
other	0.90	1.00	0.95	395
accuracy			0.98	2428
macro avg	0.97	0.90	0.92	2428

weighted avg 0.98 0.98 0.98 2428



These results are very encouraging, showing an almost excellent distinction between Greek and English and “Others” and a reliable classification of Greeklish.

1.7.5 Random Forest

Random Forest is an ensemble algorithm, which means it uses many simpler algorithms which then “vote” on a final decision. It has proven to be a good classifier on complex tasks, it combats overfitting by design (essentially by utilizing random chance in its training phase) and is still fairly easy to interpret.

The drawback is first and foremost computational, since we need to train many smaller classifiers, which may by themselves be computationally expensive (this is somewhat offset by the fact that the classifiers are independent and can be computed in parallel). Additionally, Random Forest is a non-parametric method which means that it is generally memory-intensive and may be slow to run on operational data. Finally, we also need to tune hyperparameters.

```
[48]: from sklearn.ensemble import RandomForestClassifier
```

```

forest_model = RandomForestClassifier(n_estimators=500,
                                     n_jobs=-1,
                                     criterion="entropy")
res = cross_val_res(forest_model, x_train, y_train, cv=3)
print(f"Random Forest mean macro F1: {res[0]:.4f}, std: {res[1]:.4f}")

```

Random Forest mean macro F1: 0.9292, std: 0.9111

```

[49]: forest_model = RandomForestClassifier(n_estimators=500,
                                           n_jobs=-1,
                                           criterion="entropy",
                                           verbose=1).fit(x_train, y_train)

forest_pred = forest_model.predict(x_test)
get_statistics(y_test, forest_pred)

```


```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    2.1s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed:   10.9s
[Parallel(n_jobs=-1)]: Done 434 tasks     | elapsed:   25.7s
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   30.3s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 184 tasks     | elapsed:    0.1s
[Parallel(n_jobs=8)]: Done 434 tasks     | elapsed:    0.4s
[Parallel(n_jobs=8)]: Done 500 out of 500 | elapsed:    0.4s finished

```

Macro F1: 0.9120798978330256

	precision	recall	f1-score	support
el	0.91	1.00	0.95	1012
en	0.99	1.00	0.99	942
greeklish	0.90	0.80	0.85	79
other	0.99	0.76	0.86	395
accuracy			0.95	2428
macro avg	0.95	0.89	0.91	2428
weighted avg	0.95	0.95	0.95	2428



The results are somewhat similar to our Logistic Regression classifier. The Cross-Validation score reports an equal on-average performance.

1.7.6 Adaboost

Adaboost is the logical conclusion of Random Forests, where each voter considers a very specific “rule” that needs to be followed. The next voter then considers the most important rule to distinguish between the categories for all the classes that the first could not reliably classify, and so on.

This classifier is generally more compact and competent than a simple Random Forest, but is more computationally expensive during training because we cannot train it in parallel.

```
[50]: from sklearn.ensemble import AdaBoostClassifier

ada_model = AdaBoostClassifier(n_estimators=100)
res = cross_val_res(ada_model, x_train, y_train, cv=3)
print(f"AdaBoost mean macro F1: {res[0]:.4f}, std: {res[1]:.4f}")
```


AdaBoost mean macro F1: 0.8230, std: 0.7829

```
[51]: from sklearn.ensemble import AdaBoostClassifier
```

```
ada_model = AdaBoostClassifier(n_estimators=100).fit(x_train, y_train)
ada_pred = ada_model.predict(x_test)
get_statistics(y_test, ada_pred)
```

Macro F1: 0.8173169647781657

	precision	recall	f1-score	support
el	1.00	0.87	0.93	1012
en	0.99	0.98	0.99	942
greeklish	0.80	0.42	0.55	79
other	0.68	0.98	0.80	395
accuracy			0.92	2428
macro avg	0.87	0.81	0.82	2428
weighted avg	0.94	0.92	0.92	2428



This classifier despite its complexity is even worse, being generally unable to distinguish Greekish from “Other” languages and even placing many Greek comments to the latter category.

This can be clearly seen by looking at the misclassification of adaboost compared to the misclassification of Random Forest.

```
[52]: missed = y_test != ada_pred
pd.DataFrame({"predicted": ada_pred[missed], "actual": data_test[missed].
↳labels, "text": data_test[missed].text})
```

```
[52]:
```

	predicted	actual	\	text
1373	greeklish	en		The lights look beautiful!
4411	other	el		, ...
4060	other	el		A .
2679752	other	greeklish		old expansion lordaeron server litch king
410	other	el		Π
...
4193	other	el		...
1568	other	el		A .
2577	other	el		E
4773	other	el		...
7108	other	el		

[202 rows x 3 columns]

```
[53]: missed = y_test != forest_pred
pd.DataFrame({"predicted": forest_pred[missed], "actual": data_test[missed].
↳labels, "text": data_test[missed].text})
```

```
[53]:
```

	predicted	actual	text
9602	el	other	...
9632	el	other	Syria "zgłosić broń chemiczną i podpisać konwe...
10505	el	other	Pakistańscy talibowie wyznaczają nowego przywó...
9644	el	other	...
9118	el	other	Dali-schilderij gestolen uit galerie
...
533	en	greeklish	Warwick greek voice is actor is da best
10503	el	other	2019.07.03 ...

10865	e1	other	...	
10561	e1	other		Hapana shaka rudisha .
10438	e1	other		

[118 rows x 3 columns]

1.7.7 Hyperparameter tuning

Despite the allure of Logistic Regression's properties, we will stick to the more complex model of Random Forest, since we anticipate that the sample we procured for training and testing may not necessarily be very close to the actual operational data. We thus value the stability and robustness of a non-parametric method than the computational complexity and theoretical benefits of Logistic Regression.

Since training Random Forest models is computationally intensive, we will only execute hyperparameter tuning on the most significant hyper-parameter; the number of trees which will vote during testing. This is also where our validation set comes into play.

```
[54]: estimators = []
      scores = []

      for n_estimators in tqdm([int(x) for x in np.linspace(start = 200, stop = 1200,
      num = 10)]):
          estim = RandomForestClassifier(n_estimators=n_estimators,
                                      n_jobs=-1,
                                      criterion="entropy").fit(x_train, y_train)

          score = f1_score(y_val, estim.predict(x_val), average='macro',
          zero_division=0)

          estimators.append(estim)
          scores.append(score)
```

```
100%|
  | 10/10 [07:11<00:00, 43.19s/it]
```

```
[55]: best_model = estimators[np.argmax(scores)]
      print(f"Best model {best_model} with macro F1 score of {max(scores)}")
```

```
Best model RandomForestClassifier(criterion='entropy', n_estimators=977,
n_jobs=-1) with macro F1 score of 0.9372640339591569
```

1.7.8 Annotating the operational dataset

We use our optimal classifier to identify the language of our operational dataset, which in this case are the crawled YouTube comments:

```
[56]: x_oper = vectorizer.transform(crawl_df.text)
      crawl_df["language"] = best_model.predict(x_oper)
      crawl_df
```

[56]:

					title \
0	M	N .1 () - 100	...	
1	M	N .1 () - 100	...	
2	M	N .1 () - 100	...	
3	M	N .1 () - 100	...	
4	M	N .1 () - 100	...	
...					...
2692	ΑΤΤΟΞ	ΕΧΕΙ	'0Views'(React Σ	Montages M	'0'V...
2693	ΑΤΤΟΞ	ΕΧΕΙ	'0Views'(React Σ	Montages M	'0'V...
2694	ΑΤΤΟΞ	ΕΧΕΙ	'0Views'(React Σ	Montages M	'0'V...
2695	ΑΤΤΟΞ	ΕΧΕΙ	'0Views'(React Σ	Montages M	'0'V...
2696	ΑΤΤΟΞ	ΕΧΕΙ	'0Views'(React Σ	Montages M	'0'V...
					link source \
0	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...				song
1	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...				song
2	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...				song
3	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...				song
4	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...				song
...					...
2692	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...				gaming
2693	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...				gaming
2694	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...				gaming
2695	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...				gaming
2696	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...				gaming
					text date language
0	ANTIGUAS CANCIONES DE GRECIA PAIS NATAL DE MIS...			2022-11-29	other
1	Încă o zi petrecută cu muzica voastră fantas...			2022-11-29	el
2	FelicităriSuperb \nSă fiți mereu bine \nMomen...			2022-11-29	greeklish
3	Lovely collection! Thank you <3			2022-11-29	greeklish
4	Beautiful thank you so many memories awesome ...			2022-11-29	en
...					...
2692	Copy by fantaros			2021-11-29	el
2693	Fist			2021-11-29	el
2694	Yui			2022-11-29	el
2695	Uihhii			2022-11-29	el
2696	Zzfitdt			2021-11-29	el

[2359 rows x 6 columns]

1.7.9 Exploring the operational dataset

Now that we have annotated to the best of our abilities the operational dataset, it's time to run some quick analysis on our findings.

The analysis and descriptions can be found in the `report.pdf` file. This notebook will only generate the graphs and Figures used there.

```
[57]: RESOURCE_OUTPUT = "results"

def save_plot(filename):
    path = os.path.join(RESOURCE_OUTPUT, filename)
    plt.savefig(path, bbox_inches="tight")
    print(f"Figured saved to " + path)
```

Total language frequency

```
[58]: # Define a common color palette for all graphs
palette = {"el": "blue", "en": "red", "greeklish": "green", "other": "black"}
```

```
[59]: import seaborn as sns

# I really don't have time to fix this warning, sorry :(
sns.barplot(crawl_df.language, palette=palette, legend=False)
plt.title("Post languages")
plt.xlabel("Number of observed comments")


save_plot("lang_dis.png")
plt.show()
```

C:\Users\user\AppData\Local\Temp\ipykernel_32116\3670115791.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(crawl_df.language, palette=palette, legend=False)
```

Figured saved to results\lang_dis.png



Average length of comments by language

```
[60]: languages = np.unique(crawl_df.language)

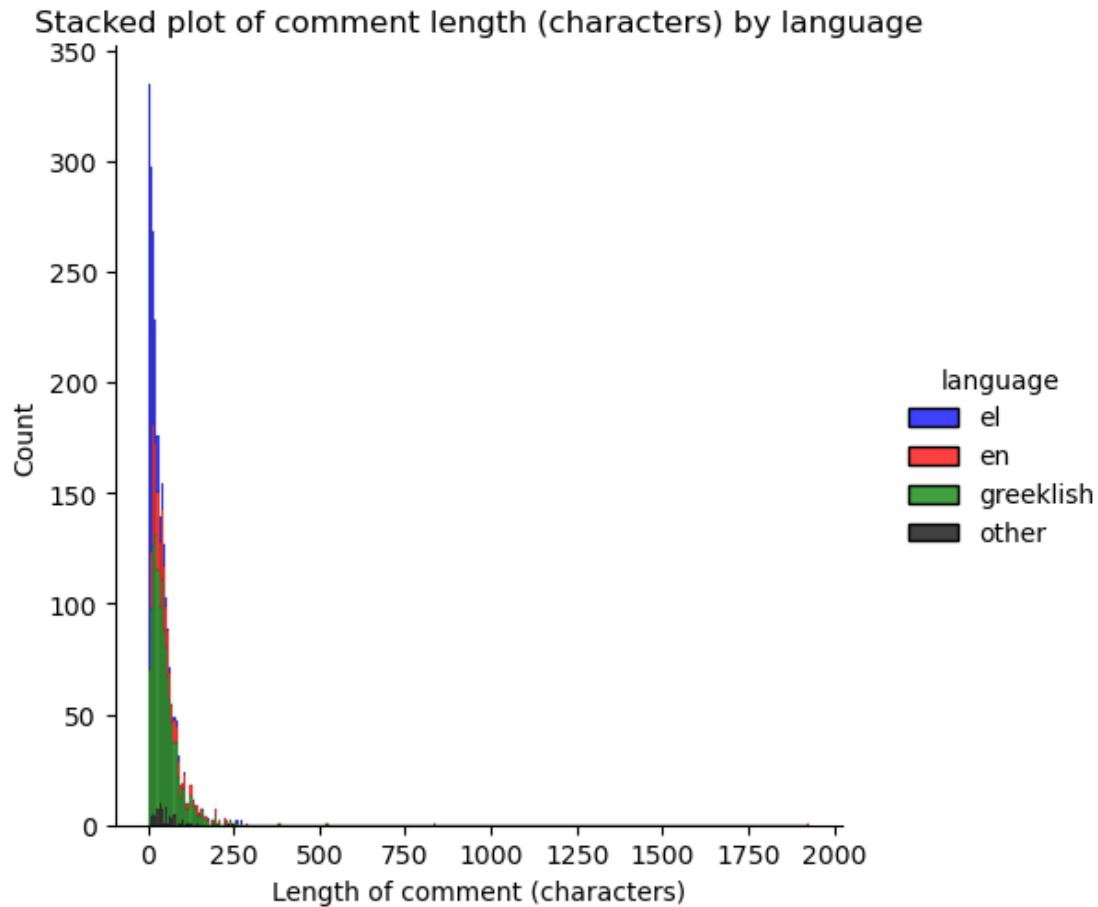
lang_col = []
len_col = []

for language in languages:
    comments_text = crawl_df.loc[crawl_df.language == language, "text"]
    comments_length = comments_text.apply(lambda x: len(x))

    lang_col += [language] * len(comments_length)
    len_col += list(comments_length)

len_df = pd.DataFrame({"language": lang_col, "comment_length": len_col})

[61]: sns.displot(len_df, x="comment_length", hue="language", multiple="stack",
    ↪ palette=palette)
plt.title("Stacked plot of comment length (characters) by language")
plt.xlabel("Length of comment (characters)")
plt.show()
```



The wild variations between the main distribution and its long tail make parsing the graph very difficult. We will thus split it in two graphs, one containing the main body of the distribution and the other the long tail.

```
[62]: fig, (ax1, ax2) = plt.subplots(1,2)
sns.histplot(len_df[len_df.comment_length<=250],
             x="comment_length",
             hue="language",
             multiple="stack",
             palette=palette,
             ax=ax1)
ax1.set_xlabel("Length of comment (characters < 250)")
ax1.set_ylabel("Number of comments")

sns.histplot(len_df[len_df.comment_length>250],
             x="comment_length",
             hue="language",
             multiple="stack",
```

```

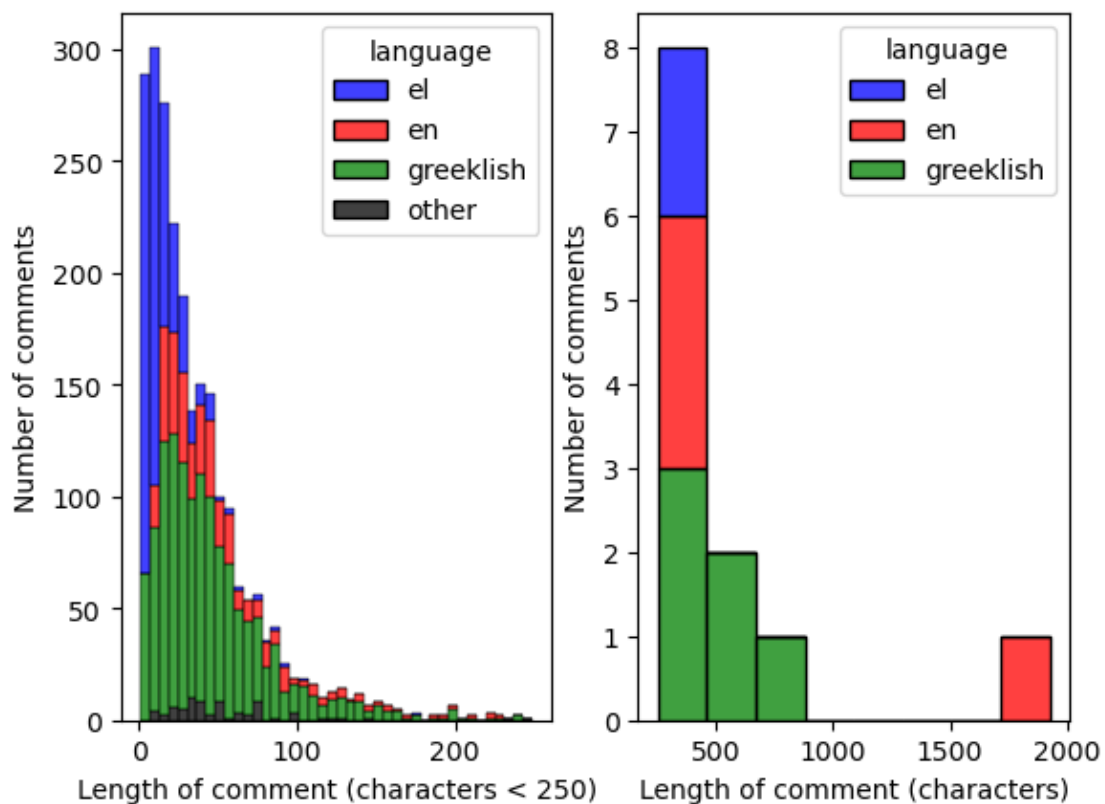
        palette=palette,
        ax=ax2)
ax2.set_xlabel("Length of comment (characters)")
ax2.set_ylabel("Number of comments")

fig.suptitle("Stacked plot of long and short comments length by language")
save_plot("length_dis.png")
plt.show()

```

Figure saved to results\length_dis.png

Stacked plot of long and short comments length by language



Emoji usage by language

```

[63]: emoji_pattern = re.compile(
    r'[\U0001F600-\U0001F64F\U0001F300-\U0001F5FF\U0001F680-\U0001F6FF'
    r'\U0001F700-\U0001F77F\U0001F780-\U0001F7FF\U0001F800-\U0001F8FF'
    r'\U0001F900-\U0001F9FF\U0001FA00-\U0001FA6F\u2600-\u26FF\u2700-\u27BF'
    r'\u2B50\u2B60\u2934\u2935\u2B05\u2194-\u2199\u21A9\u21AA\u2139\u2328'
    r'\u23CF\u23E9-\u23F3\u231A\u23F8-\u23FA\u231B\u23F0\u231A\u1F004'

```

```

        r'\u1F0CF\u1F18E\u3030\u303D]', flags=re.UNICODE
    )

    emojis_col = []

    for language in languages:
        comments_text = crawl_df.loc[crawl_df.language == language, "text"]
        comments_length = comments_text.apply(lambda x: len(emoji_pattern.
↪findall(x)))

        emojis_col += list(comments_length)

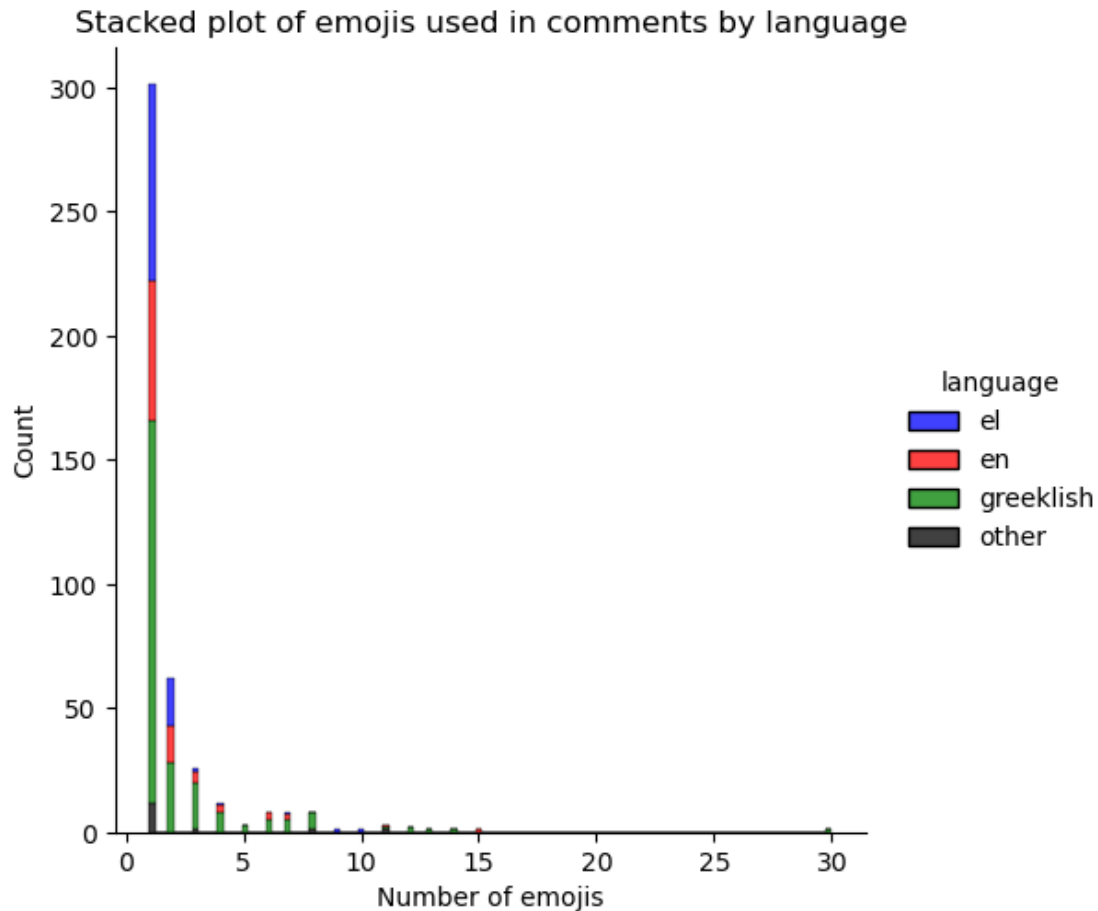
    emoji_df = pd.DataFrame({"language": lang_col, "emojis": emojis_col})

```

```

[64]: sns.displot(emoji_df[emoji_df.emojis > 0],
            x="emojis",
            hue="language",
            multiple="stack",
            palette=palette)
plt.title("Stacked plot of emojis used in comments by language")
plt.xlabel("Number of emojis")
plt.show()

```

We will perform the same operation as above for the same reasons.

```
[65]: fig, (ax1, ax2) = plt.subplots(1, 2)
sns.histplot(emoji_df[(emoji_df.emojis > 0) & (emoji_df.emojis < 10)],
             x="emojis",
             hue="language",
             multiple="stack",
             palette=palette,
             ax=ax1)
ax1.set_xlabel("Number of emojis (<10)")
ax1.set_ylabel("Number of comments")

sns.histplot(emoji_df[emoji_df.emojis > 10],
             x="emojis",
             hue="language",
             multiple="stack",
             palette=palette,
             ax=ax2)
```

```

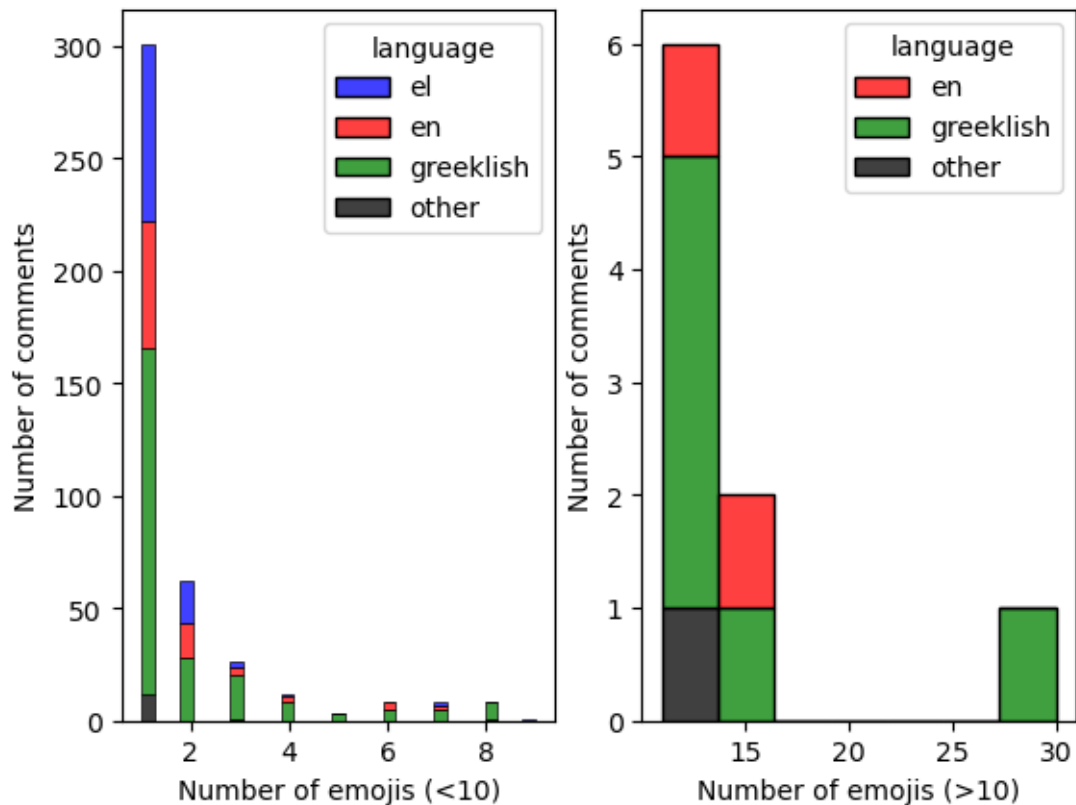
ax2.set_xlabel("Number of emojis (>10)")
ax2.set_ylabel("Number of comments")

fig.suptitle("Stacked plot of emojis used in comments by language")
save_plot("emojis_dis.png")
plt.show()

```

Figure saved to results\emojis_dis.png

Stacked plot of emojis used in comments by language



Observed language usage through time

```

[66]: date_df = crawl_df.groupby(["date", "language"]).count()
date_df

```

```

[66]:

```

	date	language	title	link	source	text
	2018-11-29	el	3	3	3	3
		en	1	1	1	1
	2019-11-29	el	49	49	49	49

	en	38	38	38	38
	greeklish	231	231	231	231
...	
2023-11-28	el	5	5	5	5
	en	3	3	3	3
	greeklish	2	2	2	2
	other	1	1	1	1
2023-11-29	greeklish	1	1	1	1

[107 rows x 4 columns]

```
[67]: date_df2 = date_df.reset_index()
date_df2.date = pd.to_datetime(date_df2.date)
date_df2.date
```

```
[67]: 0      2018-11-29
1      2018-11-29
2      2019-11-29
3      2019-11-29
4      2019-11-29
...
102     2023-11-28
103     2023-11-28
104     2023-11-28
105     2023-11-28
106     2023-11-29
Name: date, Length: 107, dtype: datetime64[ns]
```

```
[68]: import matplotlib.dates as mdates


sns.lineplot(x="date",
             y="text",
             hue="language",
             palette=palette,
             data=date_df2)

plt.title("Dates of submitted comments by language")
plt.xlabel("Year")
plt.ylabel("Number of comments")

plt.gcf().autofmt_xdate()
plt.gca().fmt_xdata = mdates.DateFormatter('%Y-%m')

save_plot("time_plot.png")
plt.show()
```

Figure saved to results\time_plot.png



1.8 Toxicity Classification

We will now attempt to build a Toxicity Classifier using an LLM. The process, prompts and decisions of this classifier can be found in the [prompts.pdf](#) report. Note that due to the project's restrictions we are only allowed to use data annotated by the LLM model.

1.8.1 Importing the data

We will import the LLM's responses and build a dataframe out of them. The annotated dataset was derived from a sample of `gold.csv`, in order to avoid feeding the model operational data.

```
[69]: with open(os.path.join("data", "chatgpt_annot.txt"), "r", encoding="utf8") as f:
      file:
          raw_annot = file.read()

      records = raw_annot.split("\n")
      len(records), records[0]
```

```
[69]: (365,
      '"E      Like      ! +1      PDT
p.s      AR      ,      AR
      !" - 1')
```

```
[70]: annotations = [record.split("-") for record in records]

comments = []
values = []
for annotation in filter(lambda x: len(x) != 1, annotations):
    comment = annotation[0]
    value = int(annotation[1])
    comments.append(comment)
    values.append(value)
```

```
[71]: toxicity_df = pd.DataFrame({"comments": comments, "toxicity": values})
toxicity_df
```

```
[71]:
```

				comments	toxicity
0	"E	Like	! +...	1	
1	"Φ	Gianuba	"	1	
2	"Congrats on your channel ,	que venha o 1 milh...			1
3	"ΤΟΣΑ ΤΕΑΕΙΑ ΒΙΝΤΕΟ ΘΑ ΚΑΝΕΙ Ο ΓΙΑΝΟΥΒΑ22"				1
4	"H	gianuba 22	pdt nu...	1	
..			
178	Ψ	raid Π	9:00, Σ ...	3	
179	M		Neltharion ...	4	
180		Σ Horde	Alliance guild;		1
181	τ	active	guild ...	1	
182	θ		realm; A ...	1	

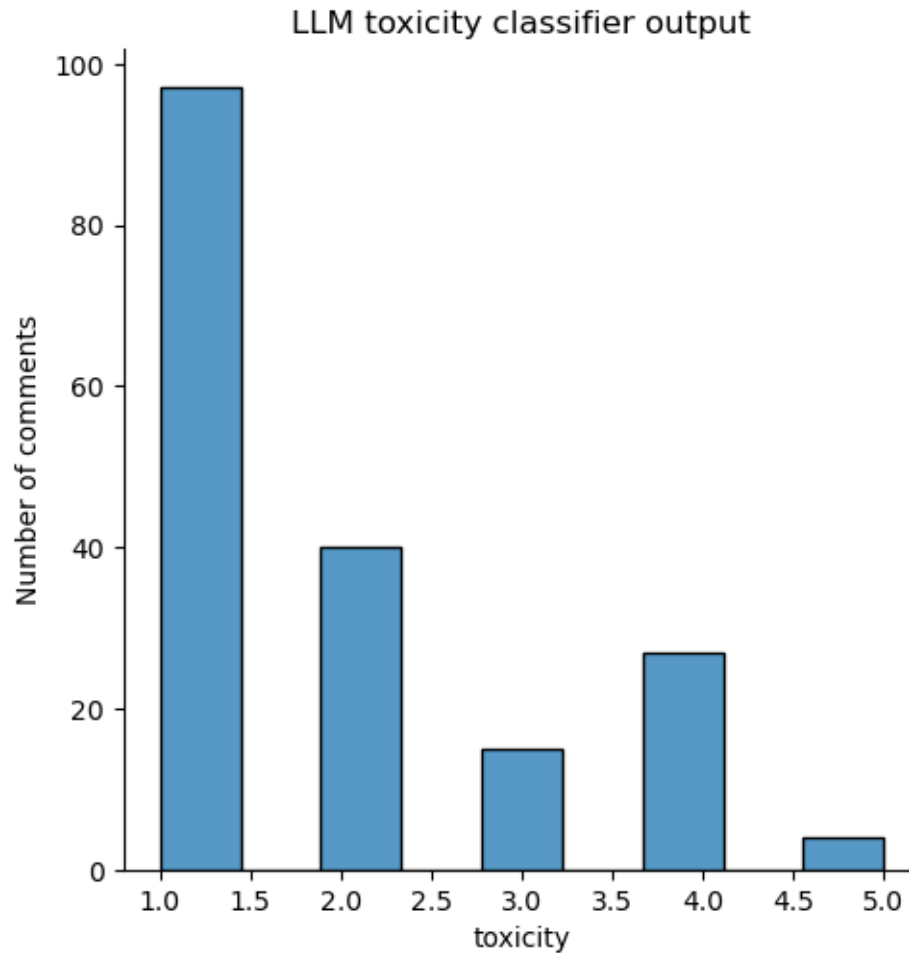
[183 rows x 2 columns]

Our dataset is comprised of 182 comments. These are of course way too few for any NLP task, but due to resource constraints and those placed by ChatGPT's frontend is the best sample we can acquire.

Our sample attempted to be as representative as possible, including data from all language categories as uniformly as possible, as well as purposefully including many toxic comments.

The overall distribution of labels produced by ChatGPT can be seen below:

```
[72]: sns.displot(toxicity_df.toxicity)
plt.title("LLM toxicity classifier output")
plt.ylabel("Number of comments")
plt.show()
```



1.8.2 Data Transformation

We will now repeat the procedure for training the same models as above. We will not be including a validation set because of the critically low amount of data we have at our disposal.

```
[73]: from sklearn.feature_extraction.text import TfidfVectorizer
      from sklearn.model_selection import train_test_split

      data_train, data_test = train_test_split(toxicity_df, random_state=42)
      # new vectorizer because of new data
      vectorizer = TfidfVectorizer().fit(data_train.comments)
      x_train = vectorizer.transform(data_train.comments)
      y_train = data_train.toxicity
      x_test = vectorizer.transform(data_test.comments)
      y_test = data_test.toxicity
```

1.8.3 Model Selection

Dummy Classifier


```
[74]: from sklearn.dummy import DummyClassifier

majority = DummyClassifier(strategy="most_frequent")
majority.fit(x_train, y_train)
majority_res = majority.predict(x_test)

get_statistics(y_test, majority_res)
```

Macro F1: 0.1408450704225352

	precision	recall	f1-score	support
1	0.54	1.00	0.70	25
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	4
4	0.00	0.00	0.00	3
5	0.00	0.00	0.00	1
accuracy			0.54	46
macro avg	0.11	0.20	0.14	46
weighted avg	0.30	0.54	0.38	46



Naive Bayes

```
[75]: naive_x_train = x_train.toarray()
naive_x_test = x_test.toarray()

naive_model = MultinomialNB()
res = cross_val_res(naive_model, naive_x_train, y_train, cv=5)
print(f"Naive Bayes mean macro F1-score {res[0]:.4f}, std: {res[1]:.4f}")
```

Naive Bayes mean macro F1-score 0.2797, std: 0.1395

C:\Users\user\anaconda3\envs\manis\Lib\site-packages\sklearn\model_selection_split.py:725: UserWarning: The least populated class in y has only 3 members, which is less than n_splits=5.
warnings.warn(


```
[76]: naive_model = MultinomialNB().fit(naive_x_train, y_train)
naive_res = naive_model.predict(naive_x_test)

get_statistics(y_test, naive_res)
```

Macro F1: 0.21913043478260869

precision	recall	f1-score	support
-----------	--------	----------	---------

1	0.55	0.96	0.70	25
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	4
4	0.50	0.33	0.40	3
5	0.00	0.00	0.00	1
accuracy			0.54	46
macro avg	0.21	0.26	0.22	46
weighted avg	0.33	0.54	0.40	46



Logistic Regression

```
[77]: with warnings.catch_warnings():
      # ignore warnings about deprecated methods in libraries
      warnings.simplefilter("ignore")

      lr = LogisticRegression(max_iter=1000)
      res = cross_val_res(lr, x_train, y_train)
      print(f"Logistic Regression mean macro F1-score {res[0]:.4f}, std: {res[1]:.4f}")
```


Logistic Regression mean macro F1-score 0.1818, std: 0.4605

```
[78]: with warnings.catch_warnings():
      warnings.simplefilter("ignore")
      lr = LogisticRegression(max_iter=1000).fit(x_train, y_train)
      lr_res = lr.predict(x_test)

      get_statistics(y_test, lr_res)
```

Macro F1: 0.21913043478260869

	precision	recall	f1-score	support
1	0.55	0.96	0.70	25
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	4
4	0.50	0.33	0.40	3
5	0.00	0.00	0.00	1
accuracy			0.54	46
macro avg	0.21	0.26	0.22	46
weighted avg	0.33	0.54	0.40	46



Random Forest

```
[79]: forest_model = RandomForestClassifier(n_estimators=200,
                                         n_jobs=-1,
                                         criterion="entropy")
res = cross_val_res(forest_model, x_train, y_train)
print(f"Random Forest mean macro F1: {res[0]:.4f}, std: {res[1]:.4f}")
```

```
C:\Users\user\anaconda3\envs\manis\Lib\site-
packages\sklearn\model_selection\_split.py:725: UserWarning: The least populated
class in y has only 3 members, which is less than n_splits=10.
  warnings.warn(
```


Random Forest mean macro F1: 0.3092, std: 0.4472

```
[80]: forest_model = RandomForestClassifier(n_estimators=200,
                                         n_jobs=-1,
                                         criterion="entropy",
                                         verbose=1).fit(x_train, y_train)
forest_pred = forest_model.predict(x_test)
get_statistics(y_test, forest_pred)
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=-1)]: Done 184 tasks    | elapsed:    0.4s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:    0.5s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 184 tasks    | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed:    0.0s finished
```

Macro F1: 0.22117647058823525

	precision	recall	f1-score	support
1	0.56	0.96	0.71	25
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	4
4	0.50	0.33	0.40	3
5	0.00	0.00	0.00	1
accuracy			0.54	46
macro avg	0.21	0.26	0.22	46
weighted avg	0.34	0.54	0.41	46



Adaboost

```
[81]: ada_model = AdaBoostClassifier(n_estimators=50)
      res = cross_val_res(ada_model, x_train, y_train, cv=3)
      print(f"AdaBoost mean macro F1: {res[0]:.4f}, std: {res[1]:.4f}")
```


AdaBoost mean macro F1: 0.2347, std: 0.2274

```
[82]: ada_model = AdaBoostClassifier(n_estimators=100).fit(x_train, y_train)
      ada_pred = ada_model.predict(x_test)
      get_statistics(y_test, ada_pred)
```

Macro F1: 0.24285714285714288

	precision	recall	f1-score	support
1	0.56	1.00	0.71	25
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	4
4	1.00	0.33	0.50	3
5	0.00	0.00	0.00	1
accuracy			0.57	46
macro avg	0.31	0.27	0.24	46

weighted avg	0.37	0.57	0.42	46
--------------	------	------	------	----



1.8.4 Classifying the crawled data

We can clearly see that none of our classifiers can make any meaningful guesses, being marginally better than a simple majority classifier, which is entirely the fault of our low amount of data. Because of the restrictions placed upon this project, we cannot augment this data in any meaningful way, and thus will continue the analysis with faulty classifiers.

```
[83]: crawl_df["toxicity"] = ada_model.predict(vectorizer.transform(crawl_df.text))
crawl_df
```

```
[83]:
```

					title \
0	M	N .1 () - 100	...	
1	M	N .1 () - 100	...	
2	M	N .1 () - 100	...	
3	M	N .1 () - 100	...	
4	M	N .1 () - 100	...	
...					...
2692	ATTOS EXEI	'OViews'(React Σ Montages M	'O'V...		

2693 ATTOE EXEI '0Views'(React Σ Montages M '0'V...
 2694 ATTOE EXEI '0Views'(React Σ Montages M '0'V...
 2695 ATTOE EXEI '0Views'(React Σ Montages M '0'V...
 2696 ATTOE EXEI '0Views'(React Σ Montages M '0'V...

	link	source \
0	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
1	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
2	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
3	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
4	https://www.youtube.com/watch?v=p5g82ta4sTk&pp...	song
...
2692	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming
2693	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming
2694	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming
2695	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming
2696	https://www.youtube.com/watch?v=6ay2HZwz2sA&pp...	gaming

	text	date \
0	ANTIGUAS CANCIONES DE GRECIA PAIS NATAL DE MIS...	2022-11-29
1	Încă o zi petrecută cu muzica voastră fantas...	2022-11-29
2	FelicităriSuperb \nSă fiți mereu bine \nMomen...	2022-11-29
3	Lovely collection! Thank you <3	2022-11-29
4	Beautiful thank you so many memories awesome ...	2022-11-29
...
2692	Copy by fantaros	2021-11-29
2693	Fist	2021-11-29
2694	Yui	2022-11-29
2695	Uihhii	2022-11-29
2696	Zzfitdt	2021-11-29

	language	toxicity
0	other	1
1	el	1
2	greeklish	1
3	greeklish	1
4	en	1
...
2692	el	1
2693	el	1
2694	el	1
2695	el	1
2696	el	1

[2359 rows x 7 columns]

1.8.5 Analysing the Toxicity of crawled data

We will now execute the scripts necessary for the analysis available at `report.pdf`.

```
[84]: def export_to_latex(df, name, col_format, caption=None):
    """
    Export a pandas DataFrame to a LaTeX file.

    :param df: The DataFrame to be exported.
    :type df: pd.DataFrame

    :param name: The name of the LaTeX file (excluding the '.tex' extension).
    :type name: str

    :param col_format: A string specifying the column formatting for the LaTeX_
    ↪table.
    :type col_format: str

    :param caption: The caption for the LaTeX table (optional).
    :type caption: str, optional

    :return: This function does not return anything.
    """
    path = os.path.join(RESOURCE_OUTPUT, name)
    df.to_latex(buf=path,
                index=False,
                formatters={"name": str.upper},
                float_format="{:.3f}".format,
                label="tab:." + name.split(".")[0],
                caption=caption,
                escape=True,
                encoding="utf-8",
                column_format=col_format)
    print(f"Dataframe exported to {path}")

def remove_emojis(text):
    return emoji_pattern.sub(r'', text)

long_col_format = '|p{10cm}|p{1cm}|'
```

Finding the most toxic language

```
[85]: toxic_lang_df = crawl_df.loc[:, ["language", "toxicity"]].groupby("language").
    ↪mean().reset_index()

caption = "Average toxicity by language."
```

```
export_to_latex(toxic_lang_df, "toxic_lang.tex", caption=caption,
               ↪col_format="|p{3.5cm}|p{1cm}|")

toxic_lang_df
```

Dataframe exported to results\toxic_lang.tex

```
[85]:      language  toxicity
0         el  1.003053
1         en  1.117949
2  greeklsh  1.004819
3        other  1.000000
```

Finding the most toxic video

```
[86]: toxic_videos_df = crawl_df.loc[:, ["link", "toxicity"]].groupby("link").mean().
       ↪sort_values("toxicity", ascending=False)
toxic_videos_df
```

```
[86]:                                     toxicity
link
https://www.youtube.com/watch?v=Iz1U4yxmRT4&pp=...  1.500000
https://www.youtube.com/watch?v=2duliv41A1I&pp=...  1.428571
https://www.youtube.com/watch?v=m7gG1w93Mq0&pp=...  1.285714
https://www.youtube.com/watch?v=d-6Y4vE3g8U&pp=...  1.250000
https://www.youtube.com/watch?v=ii2To2gvzkU&pp=...  1.166667
...
https://www.youtube.com/watch?v=KLBYBvxiTvQ&pp=...  1.000000
https://www.youtube.com/watch?v=JeNV28qWi_A&pp=...  1.000000
https://www.youtube.com/watch?v=HILSv0QV_bc&pp=...  1.000000
https://www.youtube.com/watch?v=H-TIQdWiuOg&pp=...  1.000000
https://www.youtube.com/watch?v=zHQBNoNpmog&pp=...  1.000000
```

[116 rows x 1 columns]

```
[87]: toxic_videos_df = toxic_videos_df.merge(
       crawl_df.loc[:, ["link", "title"]].drop_duplicates(),
       on="link",
       how="inner").loc[:, ["title", "toxicity"]].head(5)

toxic_videos_df.title = toxic_videos_df.title.apply(lambda x: remove_emojis(x))
caption = "The top 5 videos with the most toxic comments on average."
export_to_latex(toxic_videos_df, "toxic_videos.tex", caption=caption,
               ↪col_format=long_col_format)

toxic_videos_df
```

Dataframe exported to results\toxic_videos.tex


```
[87]:
```

		title	toxicity
0	P : A	,	1.500000
1	A E T	/ Greek Music Non...	1.428571
2	OTI BPΩ ΣTO	ORTNIT TO TPΩ CHALLNG! (ortnite G...	1.285714
3	Δ M Γ T M	(A "E : E...	1.250000
4	E	- 70 ...	1.166667

Finding videos where toxicity was uniform across time

```
[88]: # get the mean toxicity per day
toxic_time_df = crawl_df.loc[:, ["link", "date", "toxicity"]].groupby(["link", "date"]).mean().sort_values("toxicity", ascending=False)
toxic_time_df
```

```
[88]:
```

link	date	toxicity
https://www.youtube.com/watch?v=2duliv41A1I&pp=...	2022-12-29	4.000000
https://www.youtube.com/watch?v=Iz1U4yxmRT4&pp=...	2023-11-24	4.000000
https://www.youtube.com/watch?v=d-6Y4vE3g8U&pp=...	2023-06-29	2.000000
https://www.youtube.com/watch?v=q7elH0jAyEY&pp=...	2023-07-29	1.307692
https://www.youtube.com/watch?v=m7gG1w93Mq0&pp=...	2021-11-29	1.285714
...
https://www.youtube.com/watch?v=JeNV28qWi_A&pp=...	2023-11-20	1.000000
	2023-11-17	1.000000
	2023-11-15	1.000000
https://www.youtube.com/watch?v=Iz1U4yxmRT4&pp=...	2023-11-28	1.000000
https://www.youtube.com/watch?v=zHQBNoNpmog&pp=...	2023-06-29	1.000000

[327 rows x 1 columns]

```
[89]: # get the std of each link according to all dates
toxic_time_var_df = toxic_time_df.groupby(["link", "toxicity"]).std().
    ↪reset_index().sort_values("toxicity", ascending=True)
uniform_toxic_df = toxic_time_var_df[toxic_time_var_df.toxicity == 1]

# get video titles
uniform_toxic_df = uniform_toxic_df.merge(
    crawl_df.loc[:, ["link", "title"]].drop_duplicates(),
    on="link",
    how="inner").loc[:, ["title", "toxicity"]]

# remove emojis to be nice to latex
uniform_toxic_df.title = uniform_toxic_df.title.apply(lambda x:
    ↪remove_emojis(x))
caption = "Videos where comment toxicity stayed uniform over time."
# only export 30 videos (to fit in latex table)
export_to_latex(uniform_toxic_df.head(30), "toxic_uniform.tex",
    ↪caption=caption, col_format=long_col_format)
```

```
uniform_toxic_df
```

Dataframe exported to results\toxic_uniform.tex

```
[89]:
```

		title	toxicity
0	Greek Music Mix 2021 - E T Mix ...	1.0	
1	ΠΩΣ EXΑΣΑ 100€ ΣΤΟ ORTNIT! *15.000 VBUCKS* (or...		1.0
2	E - 70 ...	1.0	
3	ΚΑΘΕ KILL ΑΛΛΑΖΩ ΠΑΗΚΤΡΟΛΟΓΙΟ CHALLNG! (ortnit...		1.0
4	E E - mix	1.0	
..			
108	ΕΠΙΚΗ ΠΡΩΤΗ ΝΙΚΗ ΣΤΟ ORTNIT ft Alex (LPDudes) ...		1.0
109	ΓΙΑ ΚΑΘΕ DATH ΤΡΩΝ ΚΑΤΤΕΡΗ ΚΟΤΟΜΠΟΥΚΙΑ! (ortni...		1.0
110	N Θ - A (AI Cover)	1.0	
111	P : A ,	1.0	
112	ΕΠΕΣΤΡΕΨΑ ΣΤΟ OG ORTNIT!		1.0

[113 rows x 2 columns]

Finding videos where toxicity increases over time

```
[90]: toxic_time_incr_df = crawl_df.loc[:,["link", "date", "toxicity"]].copy()
# sort by date
toxic_time_incr_df = toxic_time_incr_df.sort_values("date")
# get lag 1 difference between dates
toxic_time_incr_df["toxicity_diff"] = toxic_time_incr_df.toxicity.diff()
# find where toxicity increases
toxic_time_incr_df = toxic_time_incr_df.loc[toxic_time_incr_df.toxicity_diff >=
↳0, :]
# get mean growth of toxicity along the dates
toxic_time_incr_df = toxic_time_incr_df.loc[:, ["link", "toxicity_diff"]].
↳groupby("link").mean()
# print most toxic videos sorted by toxicity
toxic_time_incr_df = toxic_time_incr_df.reset_index().
↳sort_values("toxicity_diff", ascending=False)
toxic_time_incr_df
```

```
[90]:
```

	link	toxicity_diff
0	https://www.youtube.com/watch?v=0sTegFKn-nQ&pp...	3.000000
1	https://www.youtube.com/watch?v=2duliv41A1I&pp...	3.000000
5	https://www.youtube.com/watch?v=Iz1U4yxmRT4&pp...	3.000000
6	https://www.youtube.com/watch?v=ZTJPZJ453dY&pp...	3.000000
7	https://www.youtube.com/watch?v=b-GnJoG6VE8&pp...	3.000000
13	https://www.youtube.com/watch?v=tGSLjD2YJCY&pp...	3.000000
14	https://www.youtube.com/watch?v=z4DMFzyCkP0&pp...	3.000000
9	https://www.youtube.com/watch?v=ii2To2gvzkU&pp...	2.500000
8	https://www.youtube.com/watch?v=d-6Y4vE3g8U&pp...	2.333333

```

2 https://www.youtube.com/watch?v=7wuh7H_PabI&pp... 2.000000
3 https://www.youtube.com/watch?v=80GFCAfVHIA&pp... 2.000000
4 https://www.youtube.com/watch?v=GHzb1liwcsI&pp... 2.000000
10 https://www.youtube.com/watch?v=ivNQq52XHPc&pp... 2.000000
11 https://www.youtube.com/watch?v=m7gG1w93Mq0&pp... 2.000000
12 https://www.youtube.com/watch?v=q7e1H0jAyEY&pp... 2.000000

```

```

[91]: toxic_time_incr_df = toxic_time_incr_df.merge(
        crawl_df.loc[:, ["link", "title"]].drop_duplicates(),
        on="link",
        how="inner").loc[:, ["title", "toxicity_diff"]]

toxic_time_incr_df.title = toxic_time_incr_df.title.apply(lambda x:
    ↪remove_emojis(x))
caption = "Videos where comment toxicity stayed increased over time."
    "The toxicity\_diff represents the average difference between
    ↪comment toxicity"
    "with lag 1 across each date."
export_to_latex(toxic_time_incr_df, "toxic_increasing.tex", caption=caption,
    ↪col_format=long_col_format)

toxic_time_incr_df

```

Dataframe exported to results\toxic_increasing.tex

```

[91]:

```

		title	toxicity_diff
0	Greek Music Mix 2021 - E T Mix ...		3.000000
1	A E T / Greek Music Non...		3.000000
2	P : A ,		3.000000
3	ΣΚΟΤΩΣΑ ΤΟΝ ΜΟΝΟΓΡΑΛ Μ 20ΒΟΜΒ !		3.000000
4	ΑΝ ΓΕΛΑΣΕΙΣ ΧΑΝΕΙΣ 500 VBUCKS! (ortnite unny M...		3.000000
5	Δ 30 PS UNRAL RANK...		3.000000
6	E - 120 (by ...		3.000000
7	E - 70 ...		2.500000
8	Δ Μ Γ Τ Μ (Α "Ε : Ε...		2.333333
9	ΝΙΚΗ ΜΟΝΟ ΜΕ ΜΥΘΙΚΑ ΟΠΛΑ CHALLNG! (ortnite Greek)		2.000000
10	ΕΒΓΑΑΑΝ ΤΟ *BUILDING* ΣΤΗΝ ΝΕΑ ΣΑΝ ΤΟΤ ΟΡΤΝΙ...		2.000000
11	Ν ΜΟΝΟ Π Challenge (ortnite OG)		2.000000
12	ΠΩΣ ΕΧΑΣΑ 100€ ΣΤΟ ΟΡΤΝΙΤ! *15.000 VBUCKS* (or...		2.000000
13	ΟΤΙ ΒΡΩ ΣΤΟ ΟΡΤΝΙΤ ΤΟ ΤΡΩΩ CHALLNG! (ortnite G...		2.000000
14	GRK 2K23 SUMMR MIX VOL. I by ΝΙΚΚΟΣ ΔΙΝΝΟ ...		2.000000

1.9 Exporting the operational dataset

```

[92]: csv_output(crawl_df, "crawl.csv")

```

File saved successfully as output\crawl.csv

Thanks for following along!

```
[93]: print(f"Notebook executed in {int((time()-start)// 60)} minutes and_  
      ↳{(time()-start) % 60:.1f} seconds")
```

Notebook executed in 51 minutes and 9.0 seconds